

Operating Systems

Instructor: Dr. Shachi Sharma

(Semester: Winter 2019)

Process Management

Process


- A program in execution
 - A process has its own address space consisting of:
 - Text region
 - Stores the code that the processor executes
 - Data region
 - Stores variables and dynamically allocated memory
 - Stack region
 - Stores instructions and local variables for active procedure calls

Process

- For every program, the OS allocates an “isolation” – unique set of code, stack, heap, data and BSS for each program.
- Each program runs as if there is only itself and none other.
- The programmer need not worry about other programs in the system.
- The OS takes care of starting and stopping programs and allocation blocks of memory to each running program
- **Virtual memory:** Each process essentially uses the same virtual address space – also known as the offset address – which equals the entire address space – X86_32: 2^{32} addresses, X86_64: 2^{64} addresses.
- The **task scheduler/process scheduler** determines each task to start and stop and everything about the process is saved in individual **Process Control Blocks (PCBs)**.
- **Each process has its own code, stack, heap, data and BSS**

Process Memory

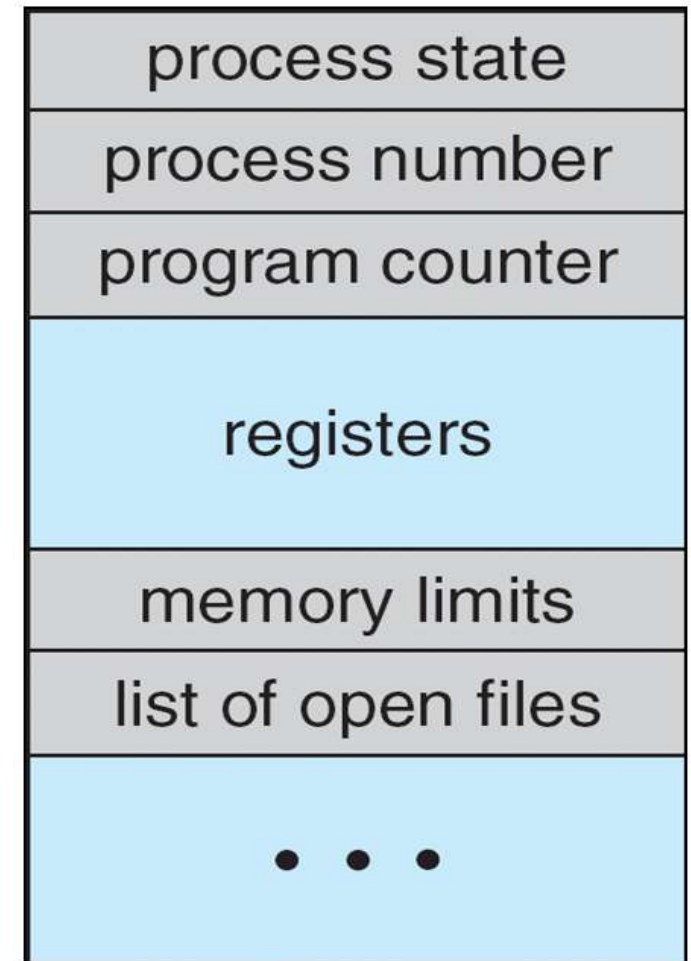
- Heap, code, data, rodata, bss – addresses increase from low to high – the registers (and subsequently the instructions) used to address it increment in the increasing order.
- Stack – address increase from higher to lower. Stack registers are decremented everytime the CPU executes a stack operation PUSH/POP/RET/CALL



Process Control Block

Information associated with each process
(also called **task control block**)

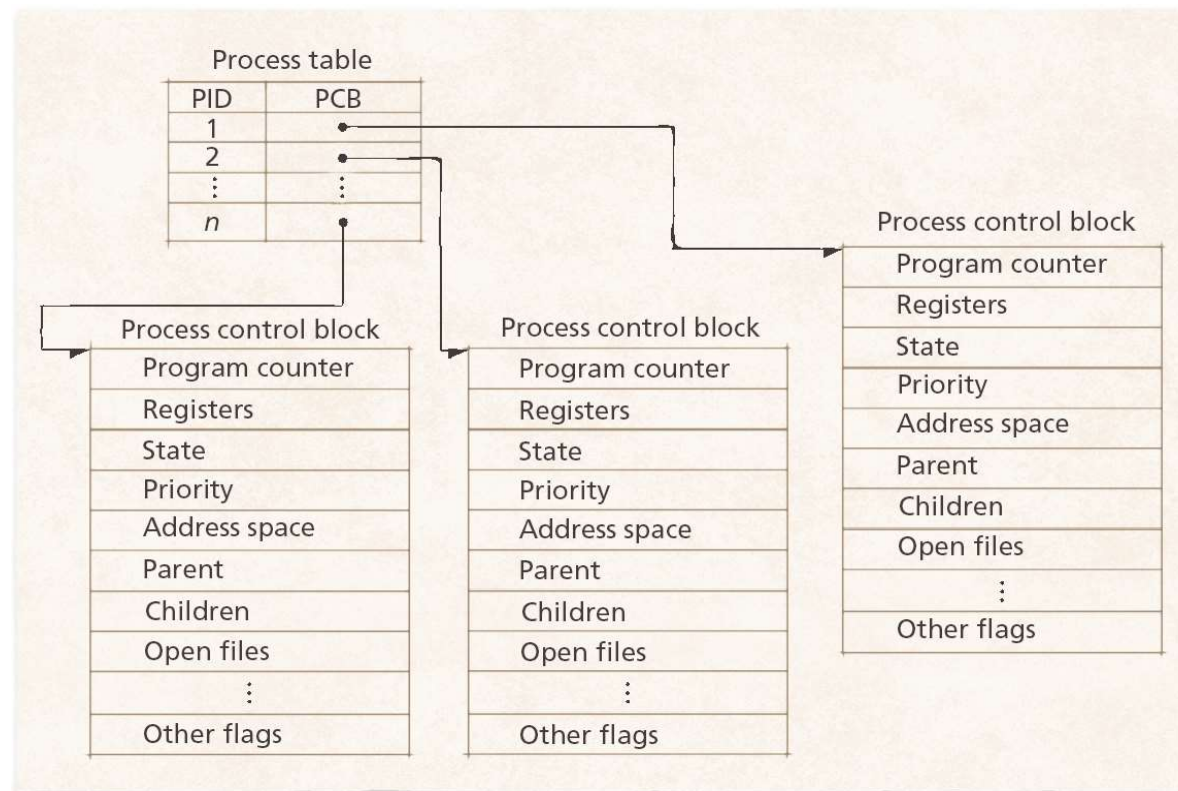
- Process state – running, waiting, etc
- *Priority (Optional)*
- Program counter – location of instruction to next execute
- CPU registers – contents of all process-centric registers
- CPU scheduling information- priorities, scheduling queue pointers
- Memory-management information – memory allocated to the process
- Accounting information – CPU used, clock time elapsed since start, time limits
- I/O status information – I/O devices allocated to process, list of open files



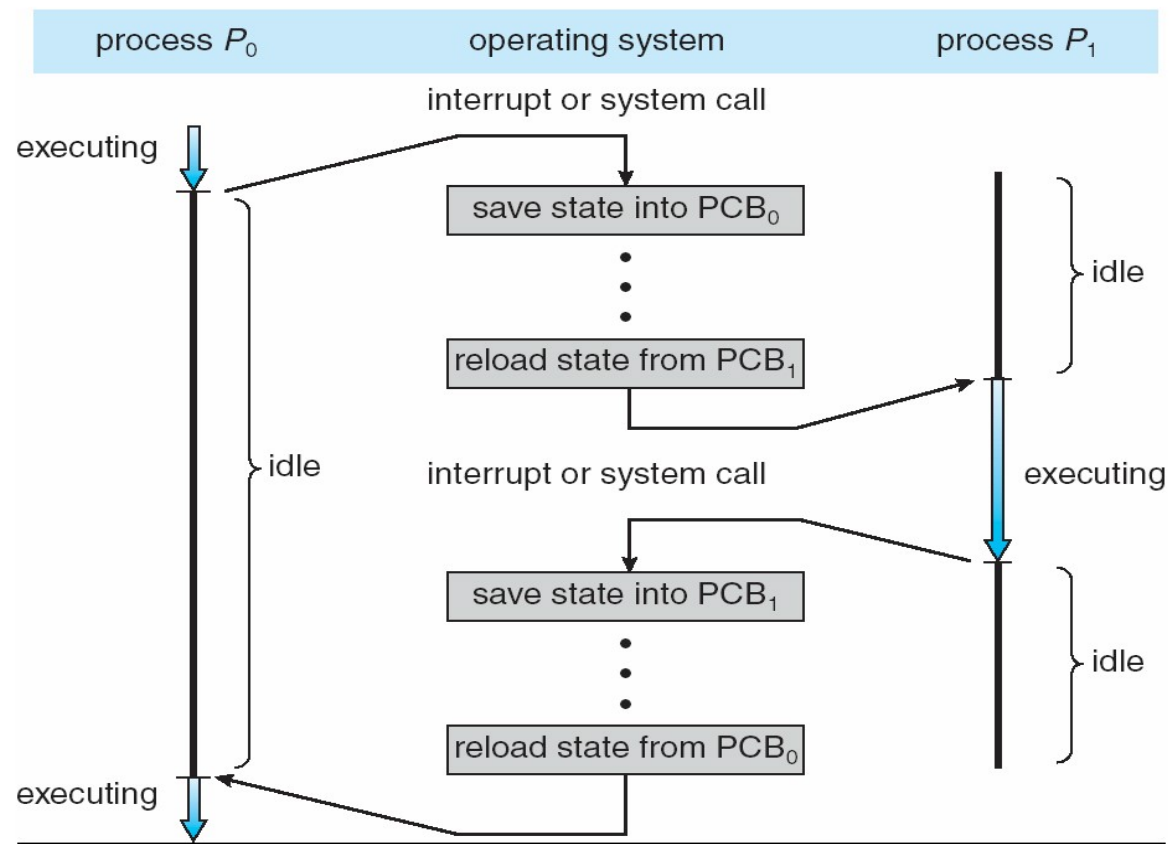
Process Table

- The OS maintains pointers to each process's PCB in a system-wide or per-user process table
- Allows for quick access to PCBs
- When a process is terminated, the OS removes the process from the process table and frees all of the process's resources

Process Table and PCBs




Task/Process Scheduler or Context Switcher



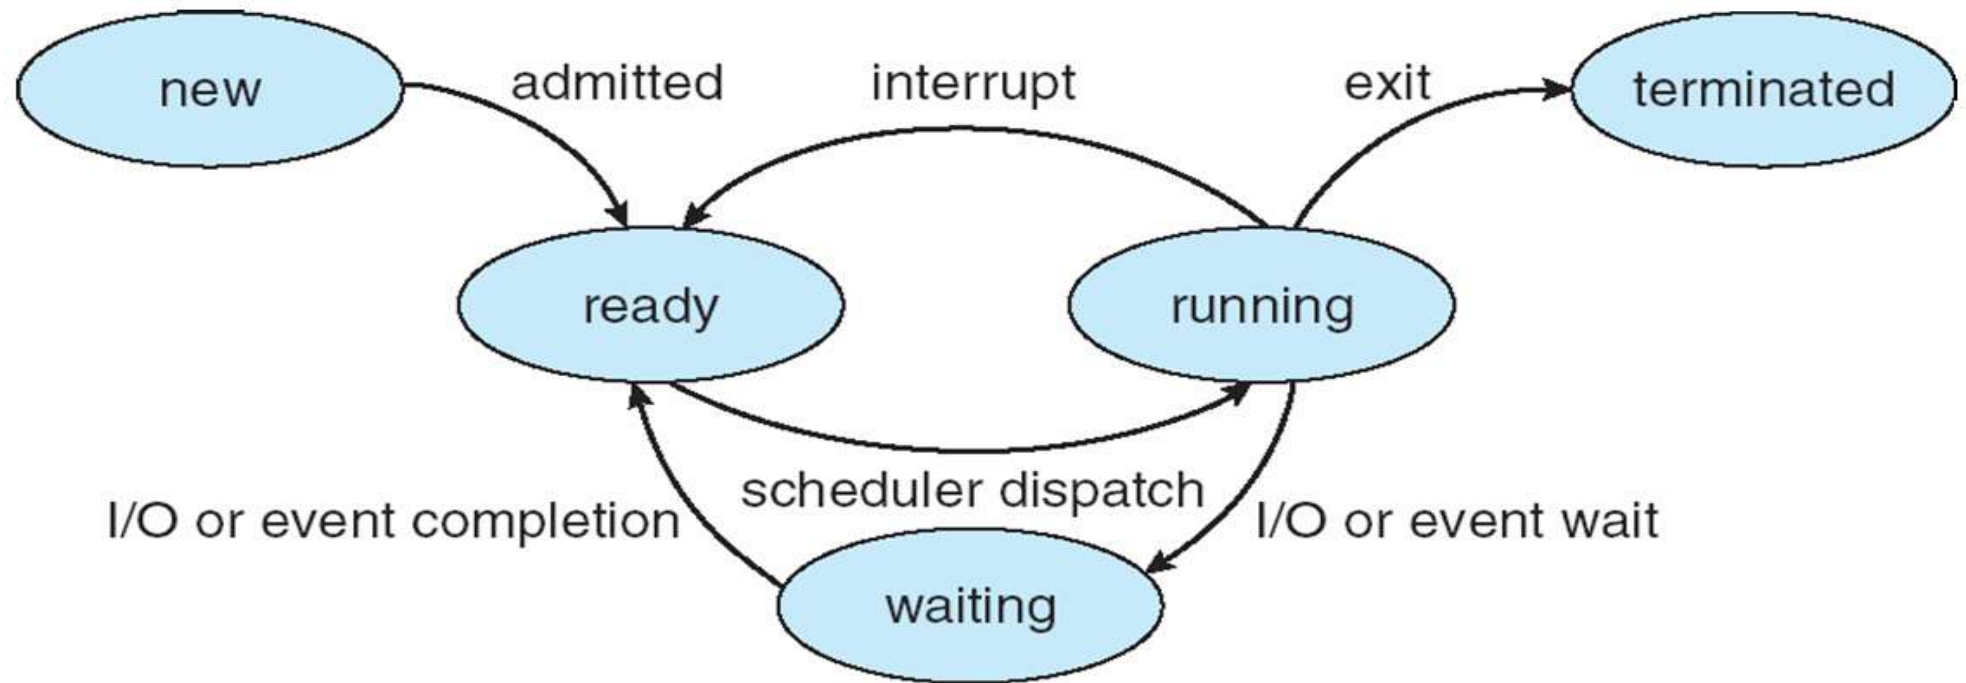
Context Switching

- Context switches
 - Performed by the OS to stop executing a *running* process and begin executing a previously *ready* process
 - Save the execution context of the *running* process to its PCB
 - Load the *ready* process's execution context from its PCB
 - Must be transparent to processes
 - Require the processor to not perform any “useful” computation
 - OS must therefore minimize context-switching time
 - Performed in hardware by some architectures

Context Switching



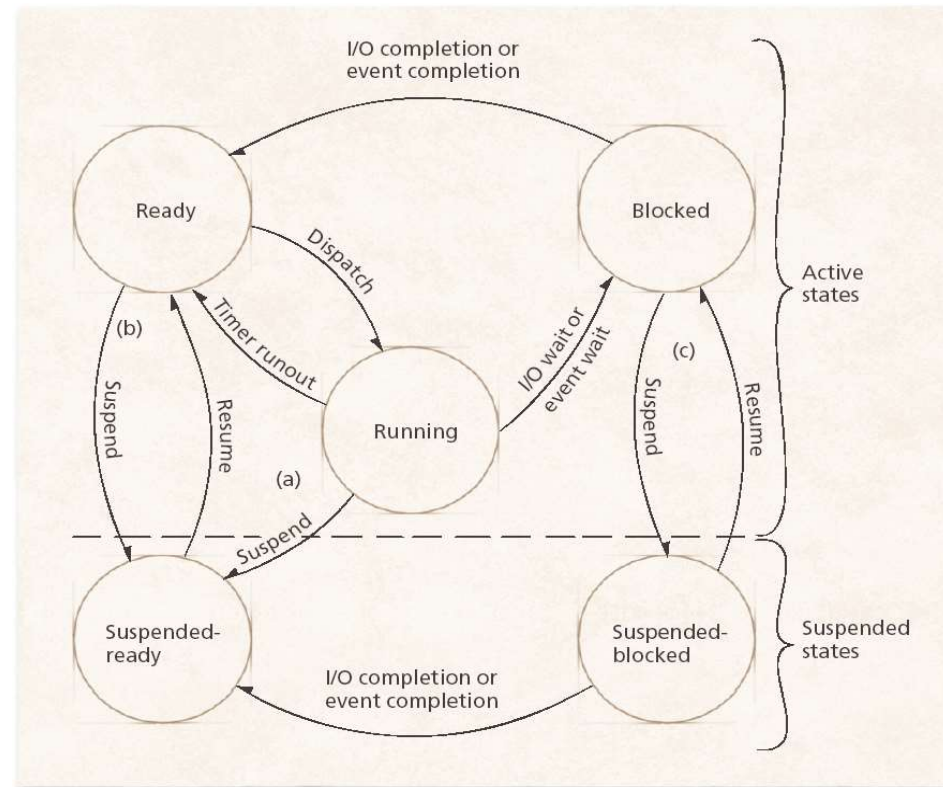
Process State



Suspend and Resume

- Suspending a process
 - Indefinitely removes it from contention for time on a processor without being destroyed
 - Useful for detecting security threats and for software debugging purposes
 - A suspension may be initiated by the process being suspended or by another process
 - A suspended process must be resumed by another process
 - Two suspended states:
 - *Suspendedready (The process is suspended and on disk but ready to run)*
 - *Suspendedblocked (The process is suspended and is waiting for an event)*

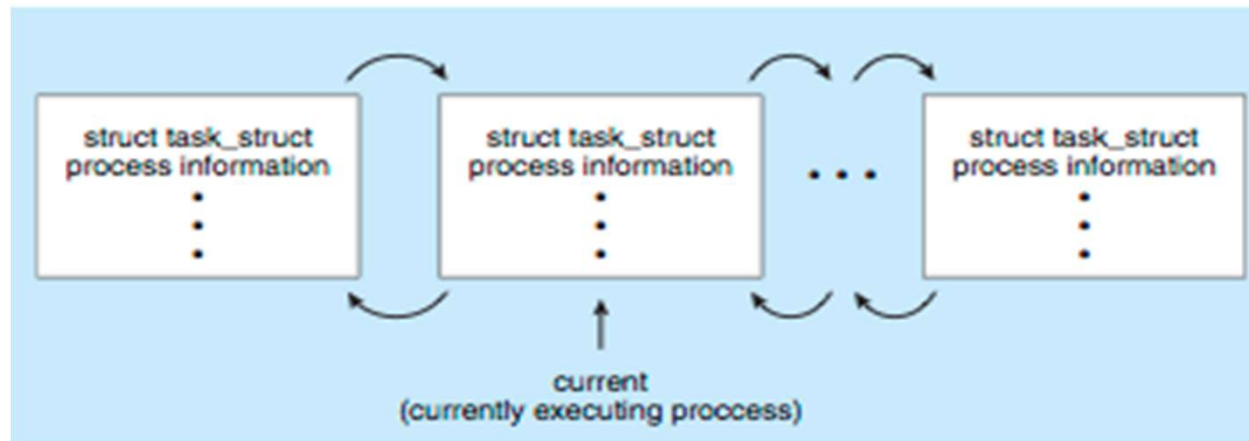
Suspend and Resume



Process Representation in Linux

- Represented by the C structure `task_struct`

```
pid_t pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice; /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```

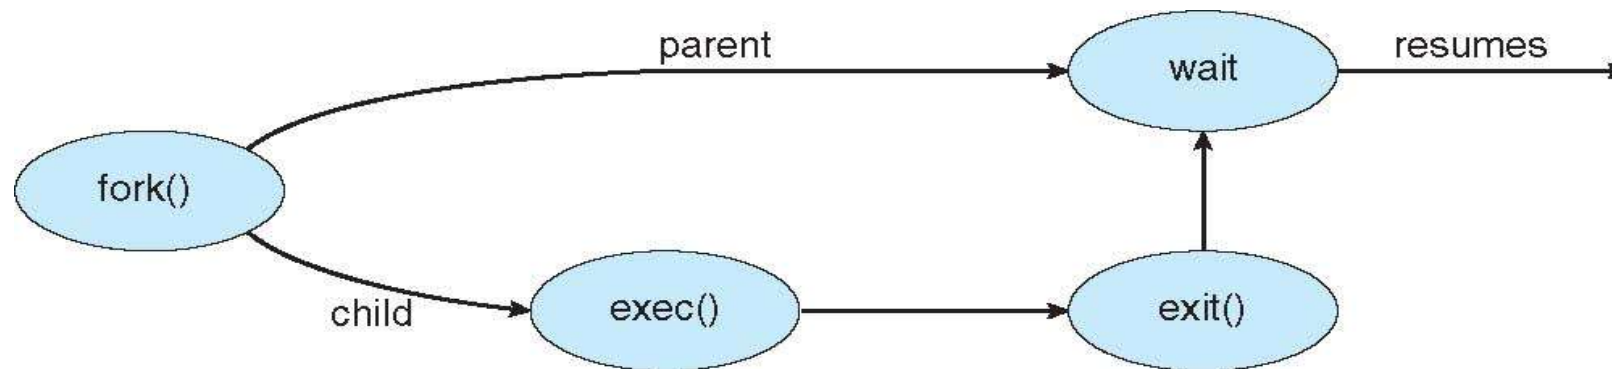


Process Creation

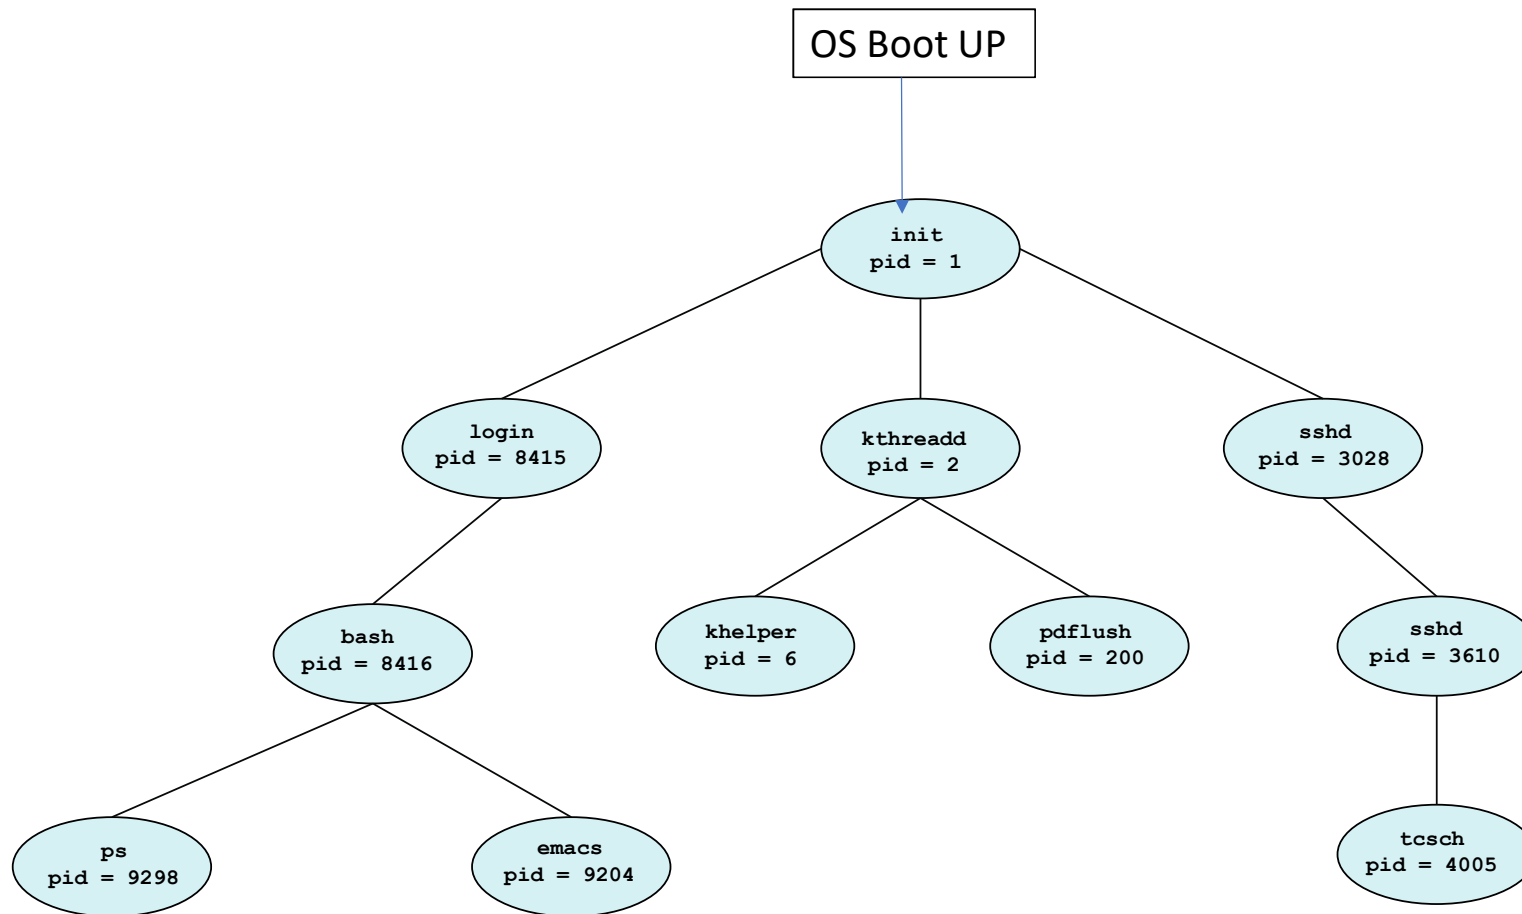
- When process is created by OS on a user request, the process is called **spawning**
- **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- Generally, process identified and managed via a **process identifier (pid)**
- Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate

Process Creation

- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - **fork()** system call creates new process
 - **exec()** system call used after a **fork()** to replace the process' memory space with a new program



Processes Tree in Linux



C Program Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

Process Termination

- Process executes last statement and asks the operating system to delete it (**exit()**)
 - Output data from child to parent (via **wait()**)
 - Process' resources are deallocated by operating system
- Parent may terminate execution of children processes (**abort()**)
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - If parent is exiting
 - Some operating systems do not allow child to continue if its parent terminates
 - All children terminated - **cascading termination**
- Wait for termination, returning the pid:

```
pid_t pid; int status;  
pid = wait(&status);
```
- If no parent waiting, then terminated process is a **zombie**
- If parent terminated, processes are **orphans**

Process Queues

